

✓ Peanut Image Denoising Autoencoder

Oleh: Crista Livia Budiman

```
import numpy as np
import pandas as pd
import random
import matplotlib.pyplot as plt
import seaborn as sns
from skimage.metrics import structural_similarity as ssim
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Conv2DTranspose, concatenate, BatchNormalization, Dropout, UpSampling2D
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
import cv2
import os
import zipfile
from glob import glob
import warnings
warnings.filterwarnings('ignore')

SEED_VALUE = 42
np.random.seed(SEED_VALUE)
random.seed(SEED_VALUE)
tf.random.set_seed(SEED_VALUE)
```

```
with zipfile.ZipFile('encoder.zip', 'r') as zip_ref:
    zip_ref.extractall('data')
```

```
DATASET_PATH = 'data'
```

```
image_folder = os.path.join(DATASET_PATH, 'B_23')
image_paths = glob(os.path.join(image_folder, '*.jpg'))
```

✓ EDA

```
shapes = []
heights = []
widths = []
channels = []
min_pixel_values = []
max_pixel_values = []

for path in image_paths:
    img = cv2.imread(path)
    if img is not None:
        shapes.append(img.shape)
        heights.append(img.shape[0])
        widths.append(img.shape[1])
        channels.append(img.shape[2])
        min_pixel_values.append(img.min())
        max_pixel_values.append(img.max())

image_info = {
    'Total Images': len(image_paths),
    'Range Height': (min(heights) if heights else None, max(heights) if heights else None),
    'Range Width': (min(widths) if widths else None, max(widths) if widths else None),
    'Channels': list(set(channels)) if channels else None,
    'Pixel Value Range': (min(min_pixel_values) if min_pixel_values else None, max(max_pixel_values) if max_pixel_values else None)
}

for key, value in image_info.items():
    print(f"{key}: {value}")
```

```
Total Images: 1074
Range Height: (600, 600)
Range Width: (600, 600)
Channels: [3]
Pixel Value Range: (np.uint8(0), np.uint8(255))
```

EDA pertama yang dilakukan adalah mengecek total gambar, shape, dan jumlah channel dari dataset.

Melalui EDA tersebut, dapat diketahui bahwa dataset berisi 1.074 gambar total, sizanya 600 x 600 pixel, memiliki 3 channel (RGB), dan range pixel nya dari rentang 0-255.

Informasi ini sangat penting karena menunjukkan bahwa semua gambar memiliki dimensi dan format yang konsisten, sehingga tidak perlu penyesuaian size saat digunakan sebagai input ke dalam model. EDA seperti ini menjadi langkah awal yang penting untuk memastikan kualitas dan kesesuaian data, sehingga proses analisis dan training model berikutnya dapat berjalan lebih akurat.

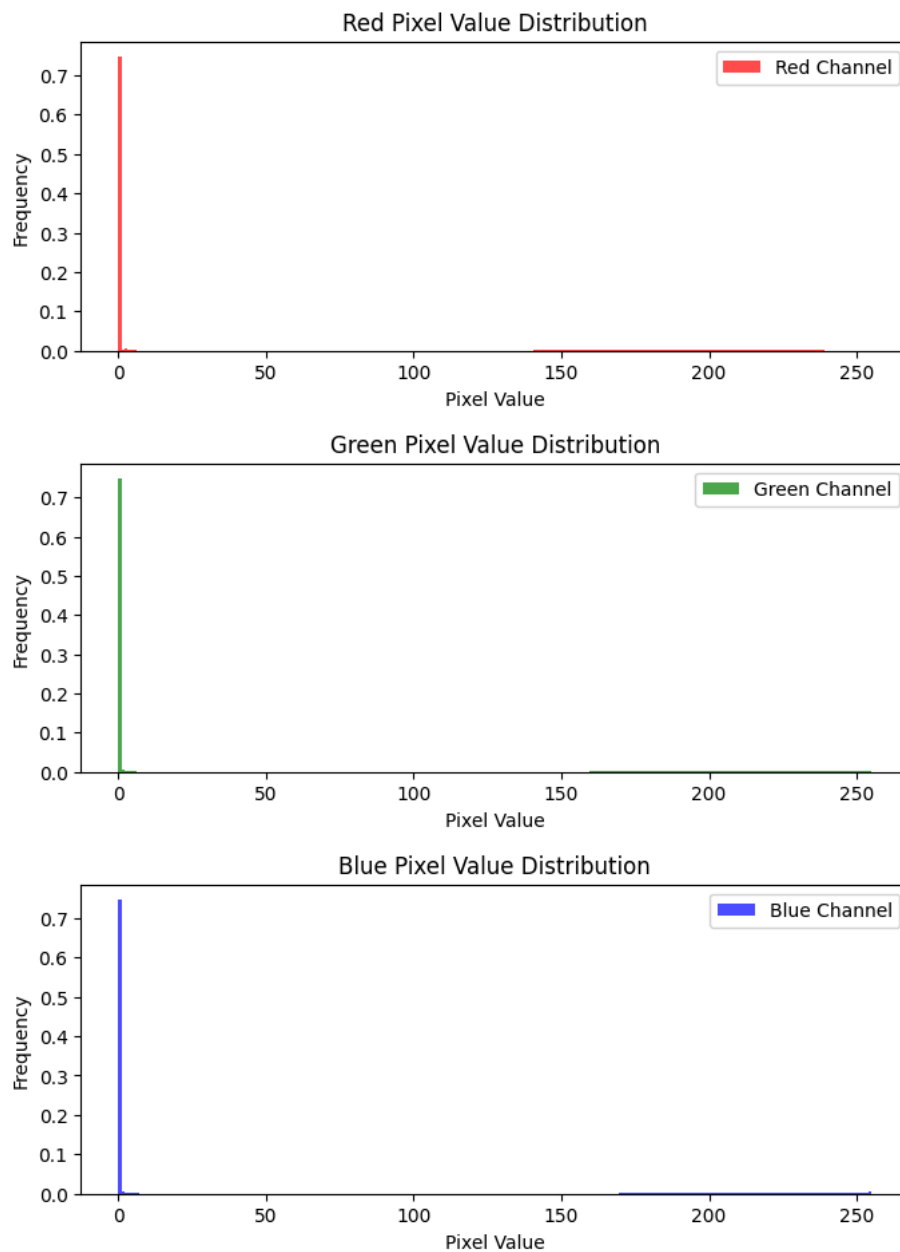
```
all_pixels = []

for path in image_paths:
    img = cv2.imread(path)
    if img is not None:
        all_pixels.append(img.reshape(-1, 3))

if all_pixels:
    all_pixels = np.concatenate(all_pixels, axis=0)

    channels = ['Red', 'Green', 'Blue']
    colors = ['red', 'green', 'blue']

    for i in range(3):
        plt.figure(figsize=(8, 3))
        plt.hist(all_pixels[:, i], bins=256, alpha=0.7, color=colors[i], label=f'{channels[i]} Channel', density=True)
        plt.title(f'{channels[i]} Pixel Value Distribution')
        plt.xlabel('Pixel Value')
        plt.ylabel('Frequency')
        plt.legend()
        plt.show()
```



EDA kedua yang dilakukan adalah mengecek distribusi nilai piksel pada channel RGB.

Melalui EDA tersebut, dapat diketahui bahwa sebagian besar nilai piksel pada ketiga channel RGB tersebut bernilai 0. Hal ini berarti mayoritas area gambar didominasi oleh warna latar belakang hitam. Terdapat sangat sedikit nilai di bawah 0.1 pada setiap nilai, yang kemungkinan besar mengindikasikan bahwa itu adalah warna objeknya.

Informasi ini penting untuk diketahui karena memberikan gambaran awal bahwa dataset memiliki rasio latar belakang yang besar dibandingkan objek utama. EDA seperti ini juga bisa membantu dalam proses segmentasi objek, terutama ketika distribusi warna memiliki perbedaan yang signifikan antara objek dan latar.

```
fig, axes = plt.subplots(2, 3, figsize=(8, 6))
for i in range(min(6, len(image_paths))):
    img = cv2.imread(image_paths[i])
    if img is not None:
        row = i // 3
        col = i % 3
        axes[row, col].imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
        axes[row, col].set_title(f'Sample {i+1}\nShape: {img.shape}')
        axes[row, col].axis('off')
plt.tight_layout()
plt.show()
```



EDA ketiga yang dilakukan adalah melihat gambar sampel beserta dimensi gambarnya.

Melalui EDA tersebut, dapat diketahui bahwa dataset berisi gambar objek kacang sebesar 600 x 600 piksel dan 3 channel warna (RGB). Setiap gambar menampilkan satu objek yang terletak di tengah dengan latar belakang hitam yang bersih dan seragam. Dari enam contoh gambar yang ditampilkan, dapat dilihat bahwa objek memiliki beberapa variasi tetapi tetap dalam posisi yang konsisten.

Informasi ini sangat penting karena menunjukkan bahwa gambar-gambar dalam dataset memiliki dimensi, komposisi, dan kualitas visual yang sama. Selain itu, posisi objek yang konsisten membuat proses training fitur oleh model menjadi lebih efektif.

✓ Preprocessing

- Split data: 80% training, 10% validasi, dan 10% testing
- Resize image: 100 x 100
- Gaussian noise: mean 0; std 0.1

✓ Splitting

```
all_images = []
for path in image_paths:
    img = cv2.imread(path)
    if img is not None:
        all_images.append(img)

all_images = [cv2.cvtColor(img, cv2.COLOR_BGR2RGB) for img in all_images]
```

```
random.shuffle(image_paths)
trainval, test = train_test_split(all_images, test_size=0.1, random_state=42)
train, val = train_test_split(trainval, test_size=1/9, random_state=42)

total = len(train) + len(val) + len(test)
print("Train:", len(train), f"({len(train)/total:.2%}")
print("Val:", len(val), f"({len(val)/total:.2%}")
print("Test:", len(test), f"({len(test)/total:.2%}")
```

```
Train: 858 (79.89%)
Val: 108 (10.06%)
Test: 108 (10.06%)
```

✓ Resize image

```
resize_shape = (100, 100)
train = np.array([cv2.resize(img, resize_shape).astype(np.float32) / 255.0 for img in train])
val = np.array([cv2.resize(img, resize_shape).astype(np.float32) / 255.0 for img in val])
test = np.array([cv2.resize(img, resize_shape).astype(np.float32) / 255.0 for img in test])
```

✓ Gaussian Noise

```
train_noise = np.clip(train + np.random.normal(0.0, 0.1, train.shape), 0., 1.)
val_noise = np.clip(val + np.random.normal(0.0, 0.1, val.shape), 0., 1.)
test_noise = np.clip(test + np.random.normal(0.0, 0.1, test.shape), 0., 1.)
```

```
train_mean = np.mean(train)
train_std = np.std(train)
val_mean = np.mean(val)
val_std = np.std(val)
test_mean = np.mean(test)
test_std = np.std(test)

train_noise_mean = np.mean(train_noise)
train_noise_std = np.std(train_noise)
val_noise_mean = np.mean(val_noise)
val_noise_std = np.std(val_noise)
test_noise_mean = np.mean(test_noise)
test_noise_std = np.std(test_noise)

print("Before Applying Gaussian Noise")
print(f"Train Mean: {train_mean:.4f}, Train Std: {train_std:.4f}")
print(f"Val Mean: {val_mean:.4f}, Val Std: {val_std:.4f}")
print(f"Test Mean: {test_mean:.4f}, Test Std: {test_std:.4f}")
print("\n")
print("After Applying Gaussian Noise")
print(f"Train Noise Mean: {train_noise_mean:.4f}, Train Noise Std: {train_noise_std:.4f}")
print(f"Val Noise Mean: {val_noise_mean:.4f}, Val Noise Std: {val_noise_std:.4f}")
print(f"Test Noise Mean: {test_noise_mean:.4f}, Test Noise Std: {test_noise_std:.4f}")
```

```
Before Applying Gaussian Noise
Train Mean: 0.1906, Train Std: 0.3453
Val Mean: 0.1886, Val Std: 0.3427
Test Mean: 0.1862, Test Std: 0.3397
```

```
After Applying Gaussian Noise
Train Noise Mean: 0.2194, Train Noise Std: 0.3327
Val Noise Mean: 0.2176, Val Noise Std: 0.3304
Test Noise Mean: 0.2152, Test Noise Std: 0.3277
```

✓ Training

```
input_shape = (100, 100, 3)
input_layer = Input(shape=input_shape, name='input_layer')
```

```
callbacks = [EarlyStopping(monitor='val_loss',patience=5,restore_best_weights=True)]
epochs = 50
batch_size = 32
```

✓ Baseline Model

```
conv2d = Conv2D(32, (3, 3), activation='relu', padding='same', name='conv2d')(input_layer)
max_pooling2d = MaxPooling2D((2, 2), padding='same', name='max_pooling2d')(conv2d)
conv2d_1 = Conv2D(64, (3, 3), activation='relu', padding='same', name='conv2d_1')(max_pooling2d)
max_pooling2d_1 = MaxPooling2D((2, 2), padding='same', name='max_pooling2d_1')(conv2d_1)
conv2d_2 = Conv2D(64, (3, 3), activation='relu', padding='same', name='conv2d_2')(max_pooling2d_1)

up_sampling2d = UpSampling2D((2, 2), name='up_sampling2d')(conv2d_2)
conv2d_3 = Conv2D(32, (3, 3), activation='relu', padding='same', name='conv2d_3')(up_sampling2d)
up_sampling2d_1 = UpSampling2D((2, 2), name='up_sampling2d_1')(conv2d_3)
conv2d_4 = Conv2D(3, (3, 3), activation='sigmoid', padding='same', name='conv2d_4')(up_sampling2d_1)
```

```
autoencoder = Model(input_layer, conv2d_4, name='autoencoder')
```

```
autoencoder.compile(
    optimizer=Adam(learning_rate=0.001),
```

```

        loss='mse',
        metrics=['mae']
    )

    autoencoder.summary()

```

Model: "autoencoder"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 100, 100, 3)	0
conv2d (Conv2D)	(None, 100, 100, 32)	896
max_pooling2d (MaxPooling2D)	(None, 50, 50, 32)	0
conv2d_1 (Conv2D)	(None, 50, 50, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 25, 25, 64)	0
conv2d_2 (Conv2D)	(None, 25, 25, 64)	36,928
up_sampling2d (UpSampling2D)	(None, 50, 50, 64)	0
conv2d_3 (Conv2D)	(None, 50, 50, 32)	18,464
up_sampling2d_1 (UpSampling2D)	(None, 100, 100, 32)	0
conv2d_4 (Conv2D)	(None, 100, 100, 3)	867

Total params: 75,651 (295.51 KB)

```

history = autoencoder.fit(
    train_noise, train,
    batch_size=batch_size,
    epochs=epochs,
    validation_data=(val_noise, val),
    callbacks=callbacks
)

```


```

Epoch 22/50
27/27 ━━━━━━━━━━━ 41s 1s/step - loss: 0.0016 - mae: 0.0136 - val_loss: 0.0016 - val_mae: 0.0134
Epoch 23/50
27/27 ━━━━━━━━━━━ 40s 1s/step - loss: 0.0016 - mae: 0.0134 - val_loss: 0.0016 - val_mae: 0.0133
Epoch 24/50
27/27 ━━━━━━━━━━━ 40s 1s/step - loss: 0.0015 - mae: 0.0132 - val_loss: 0.0015 - val_mae: 0.0132

```

```
Epoch 46/50
27/27 ----- 36s 1s/step - loss: 0.0011 - mae: 0.0110 - val_loss: 0.0011 - val_mae: 0.0108
Epoch 47/50
27/27 ----- 34s 1s/step - loss: 0.0011 - mae: 0.0110 - val_loss: 0.0011 - val_mae: 0.0108
Epoch 48/50
27/27 ----- 35s 1s/step - loss: 0.0011 - mae: 0.0109 - val_loss: 0.0011 - val_mae: 0.0108
Epoch 49/50
27/27 ----- 41s 1s/step - loss: 0.0011 - mae: 0.0109 - val_loss: 0.0011 - val_mae: 0.0106
Epoch 50/50
27/27 ----- 41s 1s/step - loss: 0.0011 - mae: 0.0108 - val_loss: 0.0011 - val_mae: 0.0106
```

```
base_train_loss = history.history['loss']
base_val_loss = history.history['val_loss']
plt.plot(base_train_loss, label='Train Loss')
plt.plot(base_val_loss, label='Val Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



Grafik menunjukkan bahwa nilai loss untuk data training dan val menurun tajam di awal, kemudian stabil mendekati nol seiring bertambahnya epoch, yang menunjukkan bahwa model baseline sudah berhasil belajar dengan baik tanpa overfitting.

✓ Modified Model

```
conv1 = Conv2D(32, (3, 3), activation='relu', padding='same')(input_layer)
conv1 = Conv2D(32, (3, 3), activation='relu', padding='same')(conv1)
pool1 = MaxPooling2D((2, 2), padding='same')(conv1)

conv2 = Conv2D(64, (3, 3), activation='relu', padding='same')(pool1)
conv2 = Conv2D(64, (3, 3), activation='relu', padding='same')(conv2)
pool2 = MaxPooling2D((2, 2), padding='same')(conv2)

# Bottleneck
conv3 = Conv2D(128, (3, 3), activation='relu', padding='same')(pool2)

# Decoder
up1 = UpSampling2D((2, 2))(conv3)
up1 = Conv2D(64, (3, 3), activation='relu', padding='same')(up1)
up1 = Add()([up1, conv2]) # skip connection

up2 = UpSampling2D((2, 2))(up1)
up2 = Conv2D(32, (3, 3), activation='relu', padding='same')(up2)
up2 = Add()([up2, conv1]) # skip connection

output_layer = Conv2D(3, (3, 3), activation='linear', padding='same')(up2)
```

```
autoencoder_modified = Model(input_layer, output_layer)

autoencoder_modified.compile(
    optimizer=Adam(learning_rate=0.001),
    loss='mse',
    metrics=['mae']
)
```

```
autoencoder_modified.summary()
```

Model: "functional"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 100, 100, 3)	0	-
conv2d (Conv2D)	(None, 100, 100, 32)	896	input_layer[0][0]
conv2d_1 (Conv2D)	(None, 100, 100, 32)	9,248	conv2d[0][0]
max_pooling2d (MaxPooling2D)	(None, 50, 50, 32)	0	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, 50, 50, 64)	18,496	max_pooling2d[0]...
conv2d_3 (Conv2D)	(None, 50, 50, 64)	36,928	conv2d_2[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 25, 25, 64)	0	conv2d_3[0][0]
conv2d_4 (Conv2D)	(None, 25, 25, 128)	73,856	max_pooling2d_1[...
up_sampling2d (UpSampling2D)	(None, 50, 50, 128)	0	conv2d_4[0][0]
conv2d_5 (Conv2D)	(None, 50, 50, 64)	73,792	up_sampling2d[0]...
add (Add)	(None, 50, 50, 64)	0	conv2d_5[0][0], conv2d_3[0][0]
up_sampling2d_1 (UpSampling2D)	(None, 100, 100, 64)	0	add[0][0]
conv2d_6 (Conv2D)	(None, 100, 100, 32)	18,464	up_sampling2d_1[...
add_1 (Add)	(None, 100, 100, 32)	0	conv2d_6[0][0], conv2d_1[0][0]
conv2d_7 (Conv2D)	(None, 100, 100, 3)	867	add_1[0][0]

```
history_modified = autoencoder_modified.fit(
    train_noise, train,
    batch_size=batch_size,
    epochs=epochs,
    validation_data=(val_noise, val),
    callbacks=callbacks
)
```

```
Epoch 1/50
27/27 ————— 115s 4s/step - loss: 0.0357 - mae: 0.1025 - val_loss: 0.0030 - val_mae: 0.0279
Epoch 2/50
27/27 ————— 112s 4s/step - loss: 0.0025 - mae: 0.0259 - val_loss: 0.0018 - val_mae: 0.0228
Epoch 3/50
27/27 ————— 113s 4s/step - loss: 0.0017 - mae: 0.0222 - val_loss: 0.0013 - val_mae: 0.0203
Epoch 4/50
27/27 ————— 140s 4s/step - loss: 0.0012 - mae: 0.0198 - val_loss: 9.8293e-04 - val_mae: 0.0181
Epoch 5/50
27/27 ————— 140s 4s/step - loss: 9.3325e-04 - mae: 0.0178 - val_loss: 8.0690e-04 - val_mae: 0.0166
Epoch 6/50
27/27 ————— 145s 4s/step - loss: 7.7321e-04 - mae: 0.0163 - val_loss: 6.9803e-04 - val_mae: 0.0154
Epoch 7/50
27/27 ————— 144s 4s/step - loss: 6.7315e-04 - mae: 0.0151 - val_loss: 6.1472e-04 - val_mae: 0.0143
Epoch 8/50
27/27 ————— 141s 4s/step - loss: 6.0199e-04 - mae: 0.0142 - val_loss: 5.4833e-04 - val_mae: 0.0134
Epoch 9/50
27/27 ————— 141s 4s/step - loss: 5.4598e-04 - mae: 0.0133 - val_loss: 5.1412e-04 - val_mae: 0.0128
Epoch 10/50
27/27 ————— 114s 4s/step - loss: 5.0042e-04 - mae: 0.0126 - val_loss: 4.6396e-04 - val_mae: 0.0119
Epoch 11/50
27/27 ————— 144s 4s/step - loss: 4.5742e-04 - mae: 0.0118 - val_loss: 4.3684e-04 - val_mae: 0.0113
Epoch 12/50
27/27 ————— 149s 5s/step - loss: 4.3207e-04 - mae: 0.0112 - val_loss: 4.1994e-04 - val_mae: 0.0108
Epoch 13/50
27/27 ————— 142s 5s/step - loss: 4.1263e-04 - mae: 0.0107 - val_loss: 4.1271e-04 - val_mae: 0.0106
Epoch 14/50
27/27 ————— 126s 5s/step - loss: 4.0042e-04 - mae: 0.0104 - val_loss: 3.8593e-04 - val_mae: 0.0100
```



```


Epoch 15/50
27/27 ----- 142s 5s/step - loss: 3.8291e-04 - mae: 0.0099 - val_loss: 3.8764e-04 - val_mae: 0.0098
Epoch 16/50
27/27 ----- 121s 5s/step - loss: 3.7604e-04 - mae: 0.0096 - val_loss: 3.6166e-04 - val_mae: 0.0093
Epoch 17/50
27/27 ----- 126s 5s/step - loss: 3.6050e-04 - mae: 0.0093 - val_loss: 3.6563e-04 - val_mae: 0.0093
Epoch 18/50
27/27 ----- 143s 5s/step - loss: 3.5652e-04 - mae: 0.0091 - val_loss: 3.5032e-04 - val_mae: 0.0090
Epoch 19/50
27/27 ----- 138s 5s/step - loss: 3.4569e-04 - mae: 0.0089 - val_loss: 3.4623e-04 - val_mae: 0.0089
Epoch 20/50
27/27 ----- 123s 5s/step - loss: 3.4053e-04 - mae: 0.0088 - val_loss: 3.4574e-04 - val_mae: 0.0089
Epoch 21/50
27/27 ----- 142s 4s/step - loss: 3.3677e-04 - mae: 0.0087 - val_loss: 3.2920e-04 - val_mae: 0.0086
Epoch 22/50
27/27 ----- 116s 4s/step - loss: 3.2758e-04 - mae: 0.0085 - val_loss: 3.3231e-04 - val_mae: 0.0086
Epoch 23/50
27/27 ----- 124s 5s/step - loss: 3.2598e-04 - mae: 0.0085 - val_loss: 3.2242e-04 - val_mae: 0.0084
Epoch 24/50
27/27 ----- 138s 4s/step - loss: 3.1922e-04 - mae: 0.0084 - val_loss: 3.2027e-04 - val_mae: 0.0084
Epoch 25/50
27/27 ----- 144s 5s/step - loss: 3.1565e-04 - mae: 0.0083 - val_loss: 3.1857e-04 - val_mae: 0.0083
Epoch 26/50
27/27 ----- 137s 4s/step - loss: 3.1256e-04 - mae: 0.0082 - val_loss: 3.1261e-04 - val_mae: 0.0082
Epoch 27/50
27/27 ----- 115s 4s/step - loss: 3.0809e-04 - mae: 0.0081 - val_loss: 3.1126e-04 - val_mae: 0.0082
Epoch 28/50
27/27 ----- 140s 4s/step - loss: 3.0561e-04 - mae: 0.0081 - val_loss: 3.0777e-04 - val_mae: 0.0081
Epoch 29/50
27/27 ----- 117s 4s/step - loss: 3.0330e-04 - mae: 0.0080 - val_loss: 3.0555e-04 - val_mae: 0.0080

```

```

modified_train_loss = history_modified.history['loss']
modified_val_loss = history_modified.history['val_loss']
plt.plot(modified_train_loss, label="Train Loss")
plt.plot(modified_val_loss, label='Val Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()

```



Grafik menunjukkan bahwa nilai loss untuk data training turun drastis, kemudian terdapat pola yang konsisten antara train loss dan val loss mendekati 0, yang mengindikasikan bahwa model ini juga belajar dengan baik dan tidak mengalami overfitting.

✓ Evaluasi

```

pred_baseline = autoencoder.predict(test_noise)
pred_modified = autoencoder_modified.predict(test_noise)

```

```

4/4 ----- 2s 332ms/step
4/4 ----- 4s 876ms/step

```

```

def calculate_ssim_list(actual, predict):
    ssim_scores = []
    for i in range(len(actual)):
        score = ssim(actual[i], predict[i], channel_axis=2, data_range=1.0)
        ssim_scores.append(score)
    return ssim_scores


```

```
ssim_baseline = calculate_ssim_list(test, pred_baseline)
ssim_modified = calculate_ssim_list(test, pred_modified)

print(f"SSIM Baseline Mean: {np.mean(ssim_baseline):.4f}")
print(f"SSIM Modifikasi Mean: {np.mean(ssim_modified):.4f}")
```

```
SSIM Baseline Mean: 0.9379
SSIM Modifikasi Mean: 0.9637
```

```
plt.figure(figsize=(12, 5))
plt.plot(ssim_baseline, label='Baseline Model')
plt.plot(ssim_modified, label='Modified Model')
plt.xlabel('Image Index')
plt.ylabel('SSIM')
plt.title('Perbandingan Grafik Pada SSIM Baseline dan Modified Model')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



Berdasarkan hasil perhitungan SSIM, diketahui bahwa Modified Model memiliki nilai rata-rata sebesar 0.9637, lebih tinggi dibandingkan dengan Baseline Model yang memperoleh nilai 0.9379. Selain itu, pada plot perbandingan SSIM, terlihat bahwa Modified Model tidak hanya memiliki nilai SSIM yang lebih tinggi, tetapi juga lebih stabil dan konsisten di berbagai sampel. Sebaliknya, Baseline Model menunjukkan fluktuasi yang cukup besar.

Oleh karena itu, dapat disimpulkan bahwa Modified Model memiliki performa yang lebih reliabel dalam merekonstruksi gambar secara menyeluruh, serta lebih akurat dan efektif dalam mempertahankan detail penting dari gambar aslinya.

Hal tersebut juga dapat dilihat pada gambar di bawah ini.

```
num_samples = 5
fig, axes = plt.subplots(4, num_samples, figsize=(15, 9))

for i in range(num_samples):
    axes[0, i].imshow(test[i])
    axes[0, i].set_title(f'Image Before Noise {i+1}')
    axes[0, i].axis('off')

    axes[1, i].imshow(test_noise[i])
    axes[1, i].set_title(f'Image After Noise {i+1}')
    axes[1, i].axis('off')

    axes[2, i].imshow(pred_baseline[i])
    axes[2, i].set_title(f'Baseline - Denoised Image {i+1}')
    axes[2, i].axis('off')

    axes[3, i].imshow(pred_modified[i])
    axes[3, i].set_title(f'Modified - Denoised Image {i+1}')
    axes[3, i].axis('off')
```

```
plt.tight_layout()
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for : nt  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for : nt  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for : nt  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for : nt  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for : nt
```

Image Before Noise 1



Image Before Noise 2



Image Before Noise 3



Image Before Noise 4



Image Before Noise 5



Image After Noise 1



Image After Noise 2



Image After Noise 3



Image After Noise 4




Image After Noise 5



Baseline - Denoised Image 1



Baseline - Denoised Image 2



Baseline - Denoised Image 3



Baseline - Denoised Image 4



Baseline - Denoised Image 5

